

Barker, Mary

Homework 2

GPU dot product of two vectors with fixed 1024 threads
using 2 methods:

- (1) Number of blocks determined by vector length
- (2) 2 blocks and iterations determined by vector length

```

/*
Mary Barker
Homework 2

Vector addition on GPU that allows for more than 1024 elements
In particular, it can do 2 different parallel memory setups:
(1) each element in the vector assigned a thread, and the number of
    blocks assigned accordingly,
(2) 2 blocks, 1024 threads each, and the algorithm iterates until
    each element in the vector has been reached

to compile: nvcc BarkerHW2.cu

OUTPUTS:
with more than 2 blocks:
Time in milliseconds= 0.0680000000000000
Last Values are A[4999] = 9998.0000000000000000 B[4999] = 4999.0000000000000000 C[4999] =
14997.0000000000000000

with only 2 blocks:
Time in milliseconds= 0.0730000000000000
Last Values are A[4999] = 9998.0000000000000000 B[4999] = 4999.0000000000000000 C[4999] =
14997.0000000000000000

*/
#include <sys/time.h>
#include <stdio.h>

//Length of vectors to be added.
#define N 5000

bool two_blocks = true;
int num_iters;
dim3 dimBlock, dimGrid;
float *A_CPU, *B_CPU, *C_CPU; //CPU pointers
float *A_GPU, *B_GPU, *C_GPU; //GPU pointers

void SetupCudaDevices()
{
    dimBlock.x = 1024;
    dimBlock.y = 1;
    dimBlock.z = 1;
    if (two_blocks)
    {
        dimGrid.x = 2;
        dimGrid.y = 1;
        dimGrid.z = 1;
        num_iters = (N - 1) / (dimGrid.x * dimBlock.x) + 1;
    }
    else
    {
        num_iters = 1;
        dimGrid.x = (N - 1) / dimBlock.x + 1;
        dimGrid.y = 1;
        dimGrid.z = 1;
    }
}

void AllocateMemory()
{
    //Allocate Device (GPU) Memory, & allocates the value of the specific pointer/array
    cudaMalloc(&A_GPU, N*sizeof(float));
    cudaMalloc(&B_GPU, N*sizeof(float));
    cudaMalloc(&C_GPU, N*sizeof(float));

```

```

        //Allocate Host (CPU) Memory
        A_CPU = (float*)malloc(N*sizeof(float));
        B_CPU = (float*)malloc(N*sizeof(float));
        C_CPU = (float*)malloc(N*sizeof(float));
    }

    //Loads values into vectors that we will add.
    void Initialize()
    {
        int i;

        for(i = 0; i < N; i++)
        {
            A_CPU[i] = (float)2*i;
            B_CPU[i] = (float)i;
        }
    }

    //Cleaning up memory after we are finished.
    void CleanUp(float *A_CPU, float *B_CPU, float *C_CPU, float *A_GPU, float *B_GPU, float *C_GPU) //
    free
    {
        free(A_CPU); free(B_CPU); free(C_CPU);
        cudaFree(A_GPU); cudaFree(B_GPU); cudaFree(C_GPU);
    }

    //This is the kernel. It is the function that will run on the GPU.
    //It adds vectors A and B then stores result in vector C
    __global__ void Addition(float *A, float *B, float *C, int n, int num_iterations_over_blocks)
    {
        int id;
        for(int i = 0; i < num_iterations_over_blocks; i++)
        {
            id = i * (blockDim.x * gridDim.x) + blockDim.x * blockIdx.x + threadIdx.x;
            if(id < n) C[id] = A[id] + B[id];
        }
    }

    int main()
    {
        int i;
        timeval start, end;

        //setup parallel structure
        SetupCudaDevices();

        //Partitioning off the memory that you will be using.
        AllocateMemory();

        //Loading up values to be added.
        Initialize();

        //Starting the timer
        gettimeofday(&start, NULL);

        //Copy Memory from CPU to GPU
        cudaMemcpyAsync(A_GPU, A_CPU, N*sizeof(float), cudaMemcpyHostToDevice);
        cudaMemcpyAsync(B_GPU, B_CPU, N*sizeof(float), cudaMemcpyHostToDevice);

        //Calling the Kernel (GPU) function.
        Addition<<<dimGrid, dimBlock>>>>(A_GPU, B_GPU, C_GPU, N, num_iters);

        //Copy Memory from GPU to CPU
        cudaMemcpyAsync(C_CPU, C_GPU, N*sizeof(float), cudaMemcpyDeviceToHost);
    }

```

```

//Stopping the timer
gettimeofday(&end, NULL);

//Calculating the total time used in the addition and converting it to milliseconds.
float time = (end.tv_sec * 1000000 + end.tv_usec) - (start.tv_sec * 1000000 + start.
    tv_usec);

//Displaying the time
printf("Time in milliseconds=%.15f\n", (time/1000.0));

// Displaying vector info you will want to comment out the vector print line when your
//vector becomes big. This is just to make sure everything is running correctly.
for(i = 0; i < N; i++)
{
    //printf("A[%d] = %.15f B[%d] = %.15f C[%d] = %.15f\n", i, A_CPU[i], i, B_CPU[i],
        i, C_CPU[i]);
}

//Displaying the last value of the addition for a check when all vector display has been
commented out.
printf("Last Values are A[%d]=%.15f B[%d]=%.15f C[%d]=%.15f\n", N-1, A_CPU[N-1],
    N-1, B_CPU[N-1], N-1, C_CPU[N-1]);

//You're done so cleanup your mess.
Cleanup(A_CPU, B_CPU, C_CPU, A_GPU, B_GPU, C_GPU);

return(0);
}

```